

LLM'S

FIRST WEEK

Gülşen Sabak

Bogazici University Computer Engineering

AGENDA

- What is an LLM?
- How does it differ from other types of ML models?
- Transformer Architecture
- Why is this revolutionary for nlp?
- Pre-training of LLM's
- Fine tuning on LLM's
- The difference between pre-training and fine-tuning
- Tokenization on LLM's
- Embeddings on LLM's
- Evaluation of LLM's
- Real World Applications of LLM's

WHAT IS AN LLM?


- base concept -> next word prediction
- neural networks specifically designed to process and generate human language at scale.
- They learn patterns in language by training on vast amounts of text data, developing an ability to understand context and generate coherent responses across a wide range of topics.

HOW DOES IT DIFFER FROM OTHER TYPES OF ML MODELS?

- **traditional ml models:**
 - handle structured data with clear input-output relationship
 - excel at specific tasks
- **earlier NLP models:**
 - often specialized for specific tasks and required separate models for different functions
 - use simpler architecture and smaller datasets
 - focusing on specific linguistic features or patterns
- **LLM's:**
 - handle unstructured text data
 - perform multiple tasks without being explicitly programmed for each one
 - use sophisticated transformer architecture
 - process long range dependencies in text and understand context


TRANSFORMER ARCHITECTURE

- The Transformer architecture is a foundational model in natural language processing (NLP) and serves as the backbone for large language models (LLMs). Introduced in the paper "Attention is All You Need"
- The core innovation of Transformers is their ability to process all words in a sequence simultaneously while understanding how each word relates to all others, primarily through the self-attention mechanism.



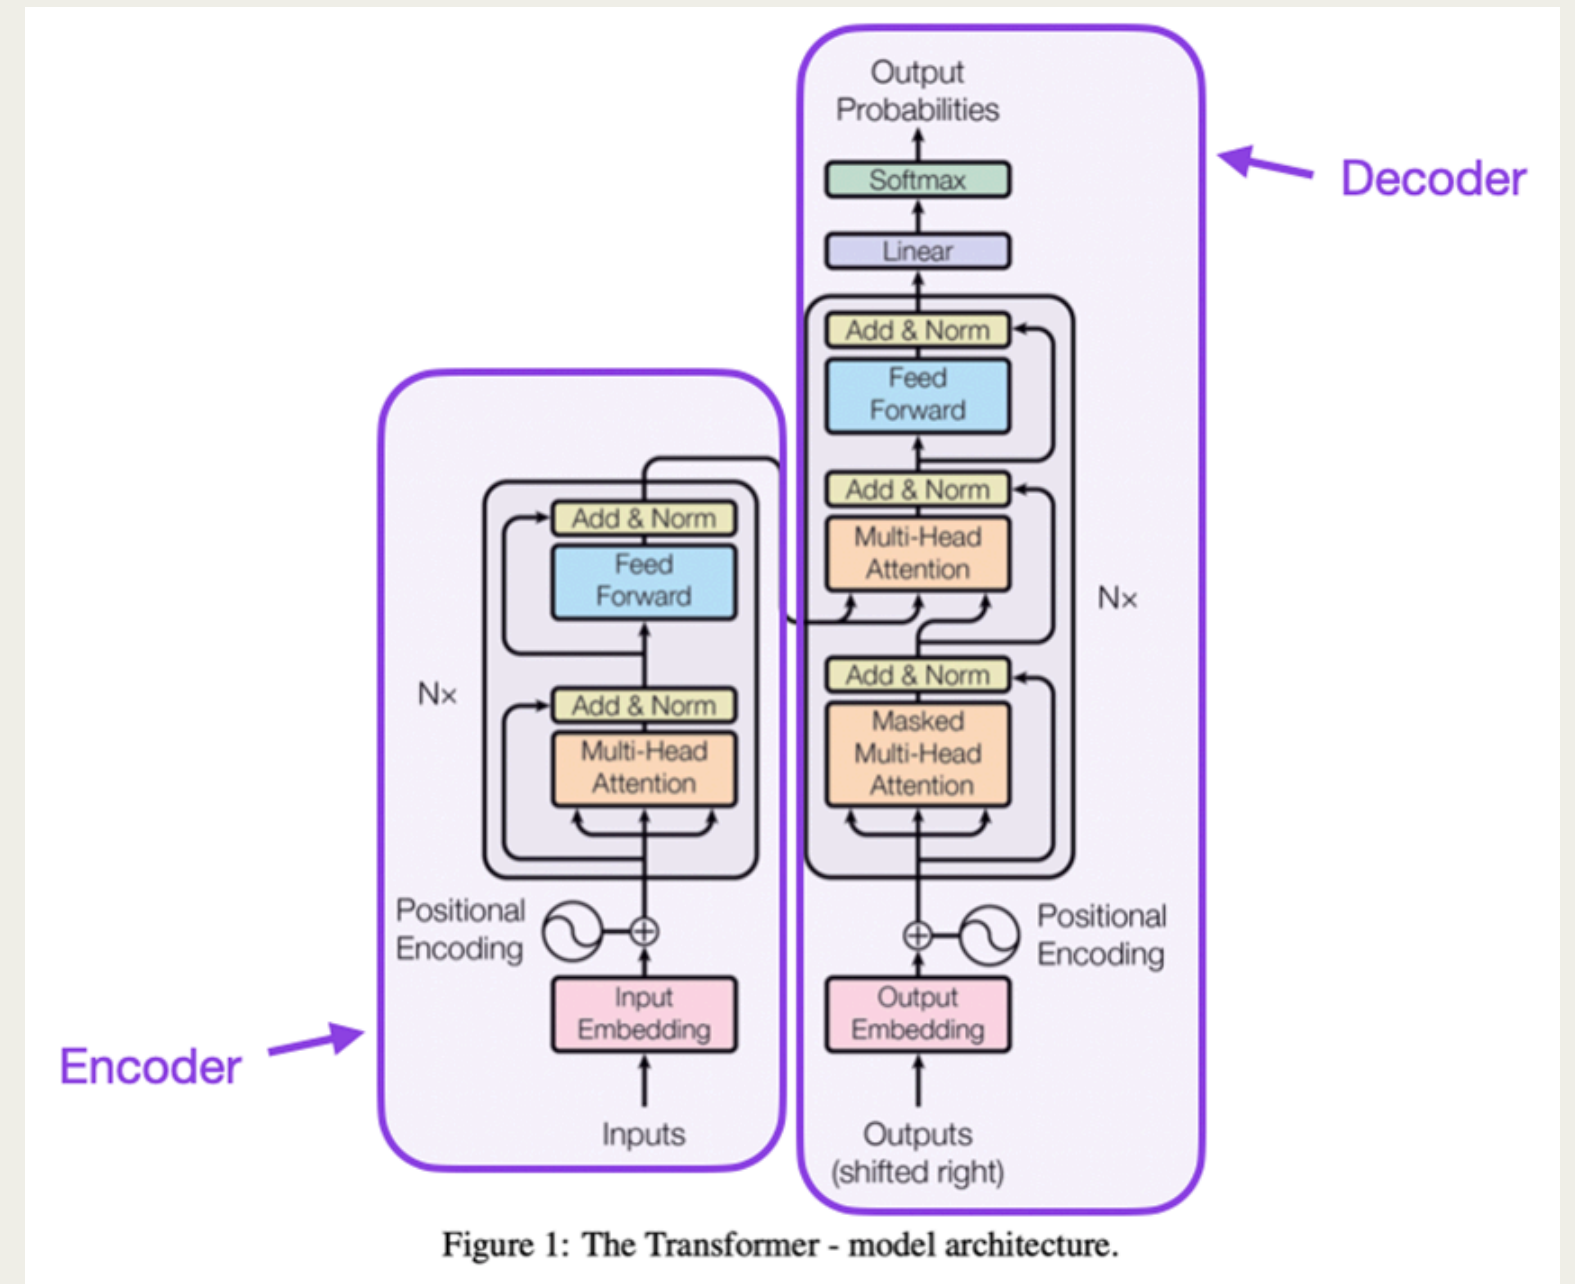
TRANSFORMER ARCHITECTURE

- **Embedding Layer:**
 - Trainable vector embedding space. Each token represented as a vector and occupies a unique location within that space.
- **Positional Encoding:**
 - Since Transformers process all words simultaneously, they need a way to understand word order
 - Sinusoidal position embeddings are added to word embeddings to encode position information
- Token embedding + positional encoding works parallel for each token



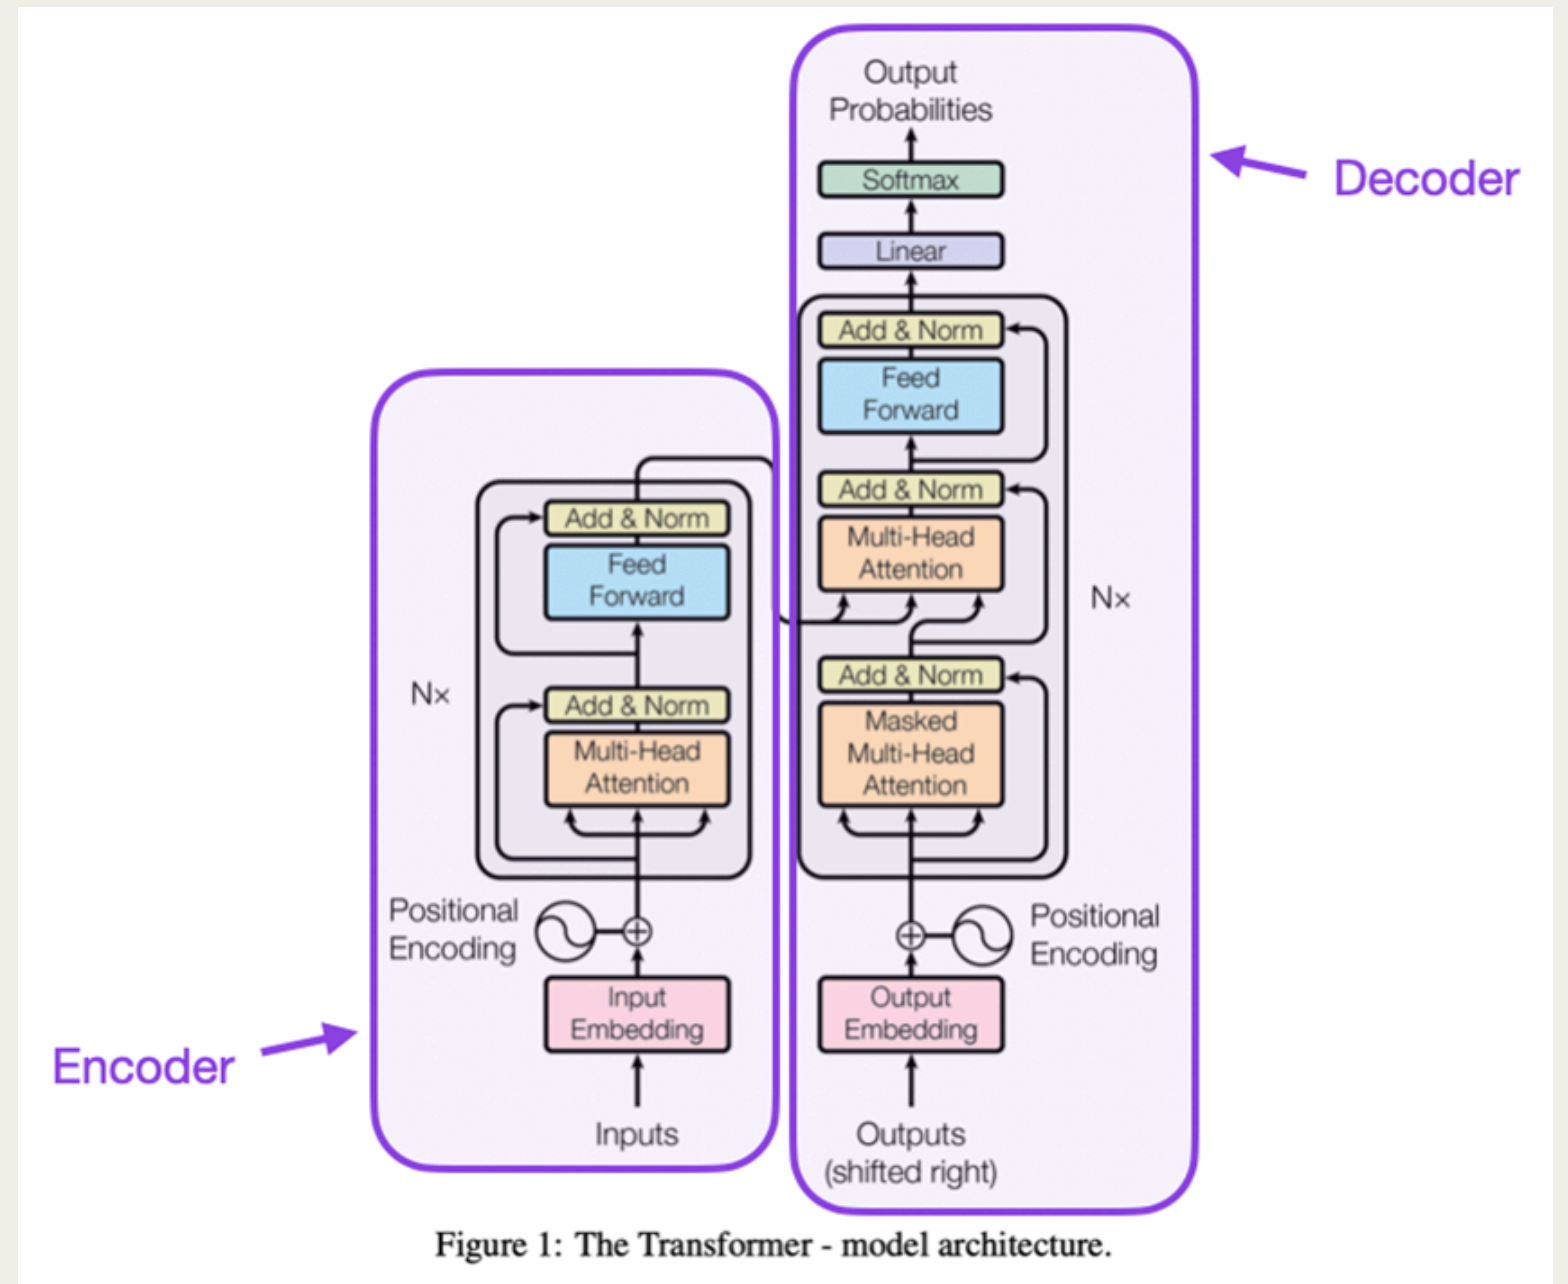
TRANSFORMER ARCHITECTURE

- **Self Attention Layer:**
 - The heart of the Transformer that allows it to weigh the importance of different words in relation to each other
 - For each word, it creates three vectors: Query (Q), Key (K), and Value (V)
 - The model computes attention scores by comparing each word's query vector with every other word's key vector
 - These scores determine how much focus to place on other words when encoding the current word
 - Captures the contextual dependencies between words



TRANSFORMER ARCHITECTURE

- **Multi-Head Attention:**
 - Multiple sets of attention mechanisms run in parallel, independently
 - Each "head" can focus on different types of relationships
 - Example: One head might focus on subject-verb relationships, while another captures adjective-noun relationships
 - 12 – 100 layers are enough



TRANSFORMER ARCHITECTURE

- **Feed Forward:**
 - Process the output to learn complex relationships
 - helps to create more meaningful and complex embeddings
 - with the help of activation functions like ReLU, model is being stronger
 - It increases the data dimension(x4) and then reduce it. With that way model learns the hidden representations

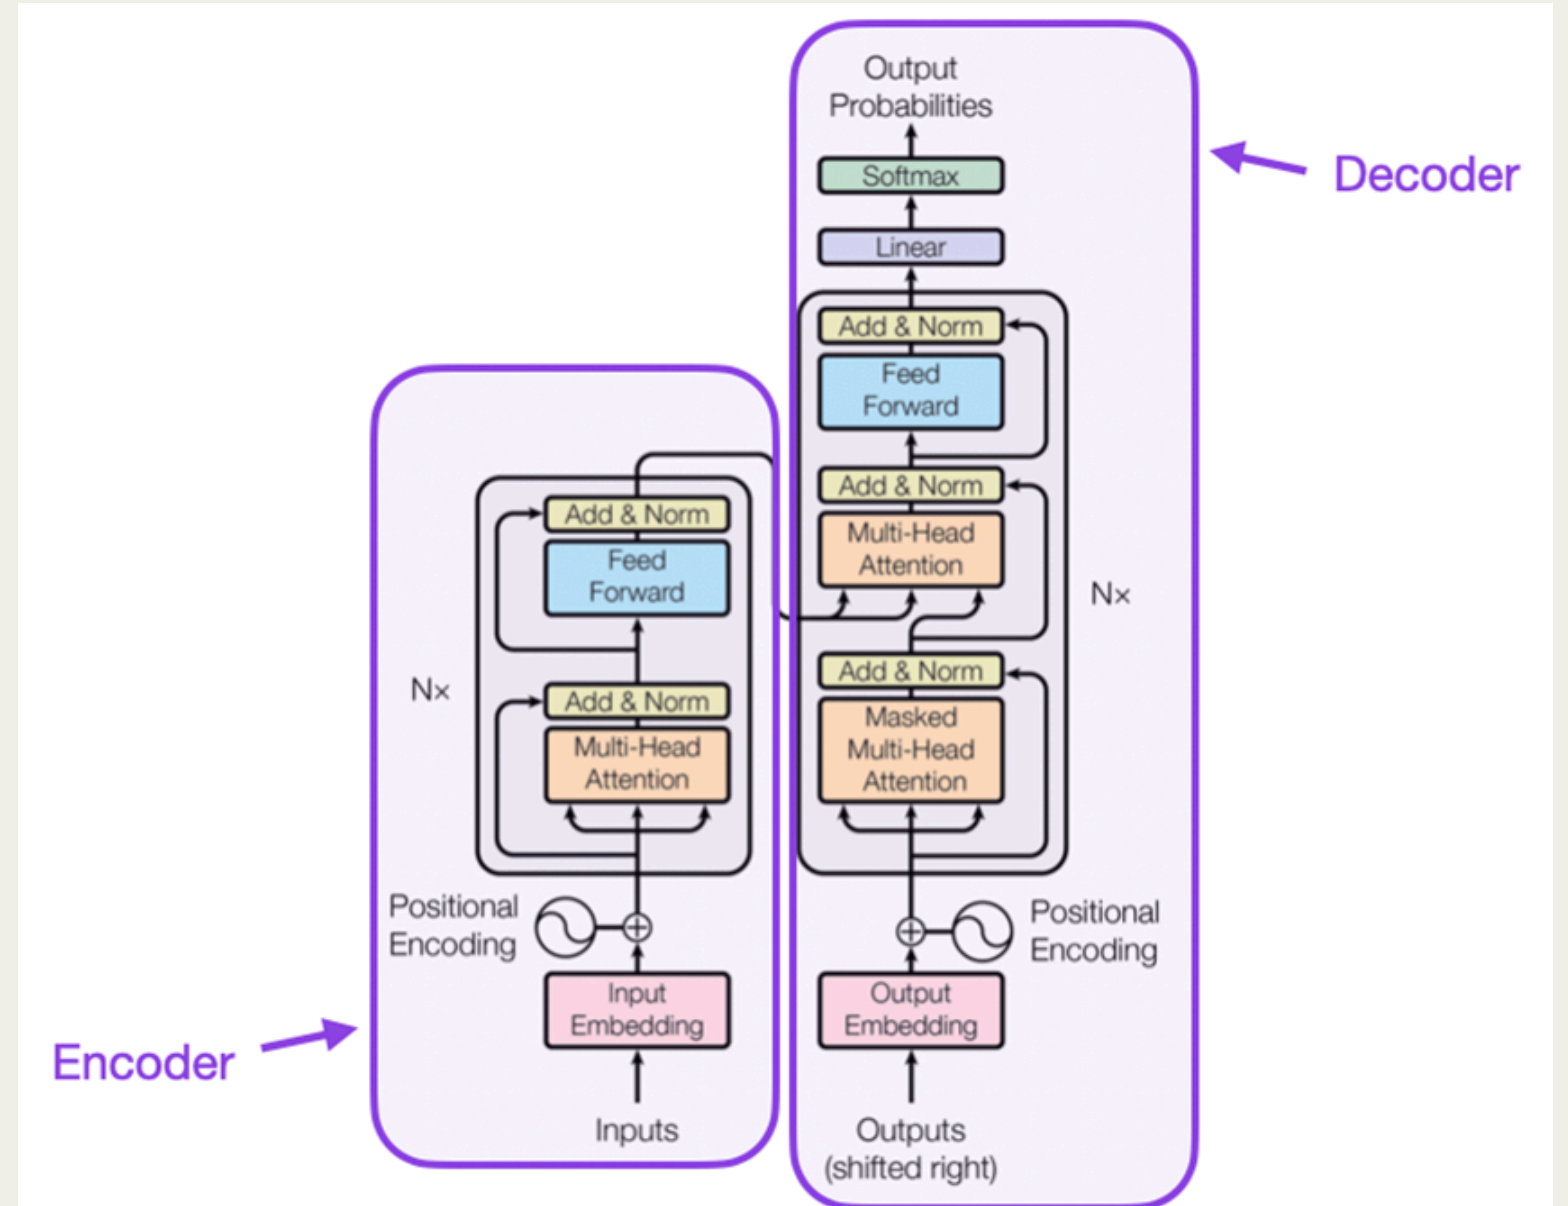


Figure 1: The Transformer - model architecture.

WHY IS THIS REVOLUTIONARY FOR NLP?

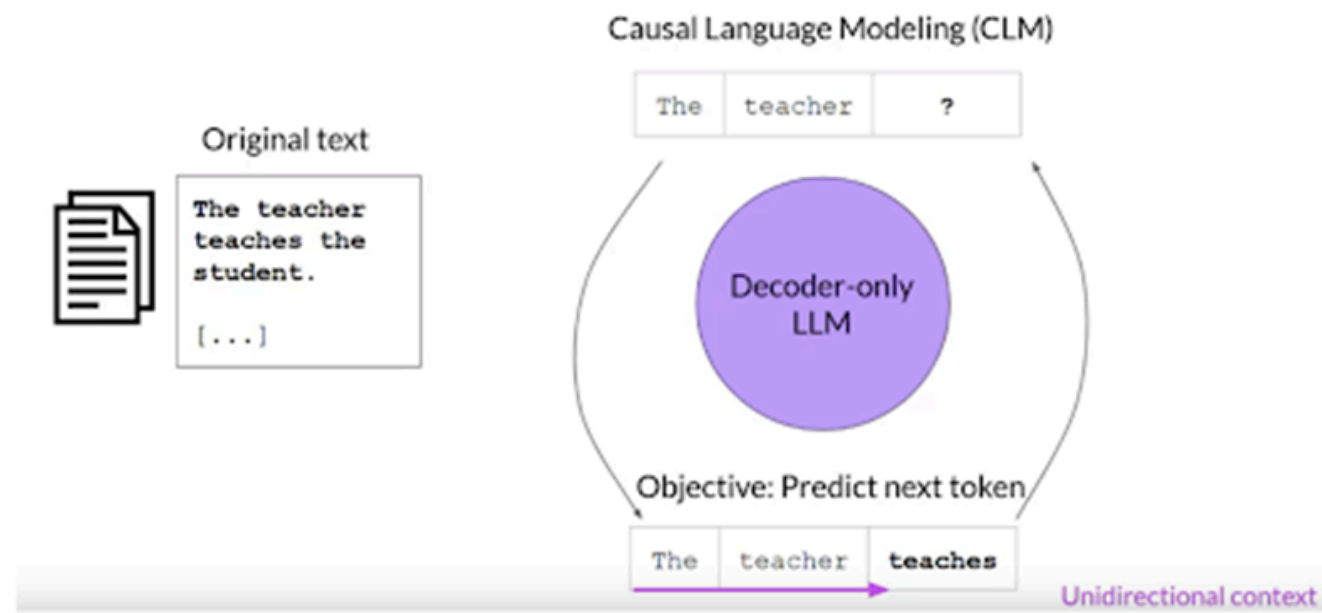
- **Parallel Processing:**
 - Unlike RNNs, which process text sequentially, Transformers process all words simultaneously
 - This enables much faster training and inference
 - Makes it practical to train on massive datasets
- **Better Long-Range Dependencies:**
 - Can directly model relationships between any two words, regardless of distance
 - Previous architectures struggled with long-distance relationships
- **Versatility:**
 - The same basic architecture works well for many different NLP tasks
 - Can be adapted for text generation, translation, summarization, and more
 - Supports transfer learning effectively

PRE TRAINING OF LLM'S

- LLM's encode a deep statistical representation of language. Understanding happens during pre-training
- Model is pre-trained with unstructured textual data
- During pre-training model weights get updated to minimize loss
- Steps of pre-training:
 - Data Collection and Preparation
 - Gathering massive amounts of text data from diverse sources (books, websites, articles)
 - Tokenizing text into subword units that the model can process
 - Objective Setting
 - Most common is the "next token prediction" task (autoregressive training)
 - The model learns to predict the next word given previous context
 - Some models also use masked language modeling, predicting hidden words in a sentence

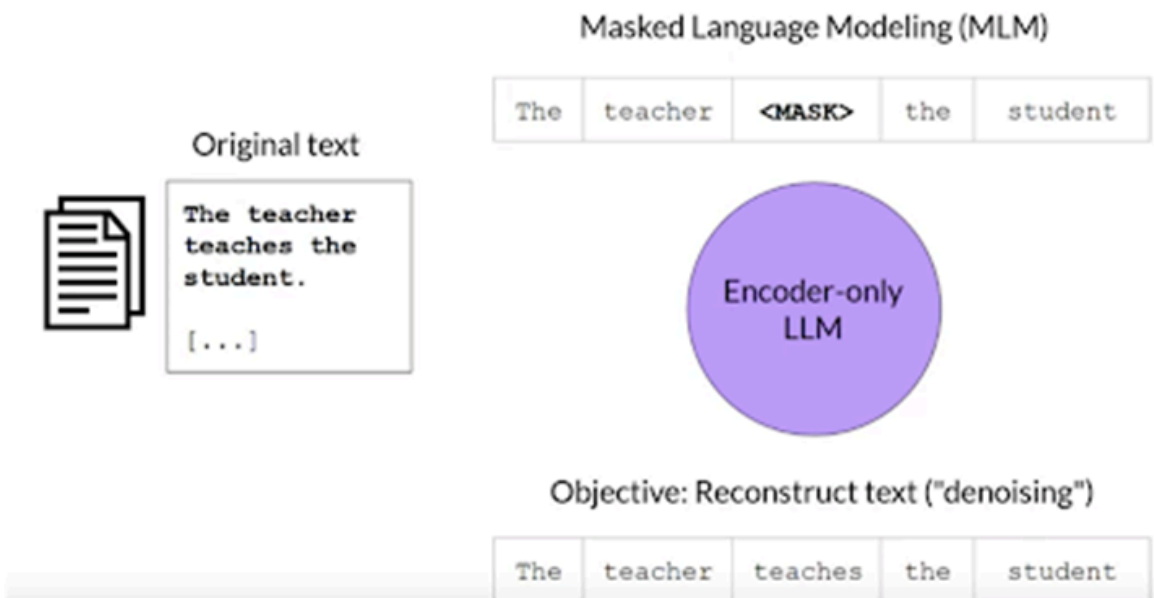
PRE TRAINING OF LLM'S

Pre-training of Decoder-only LLM



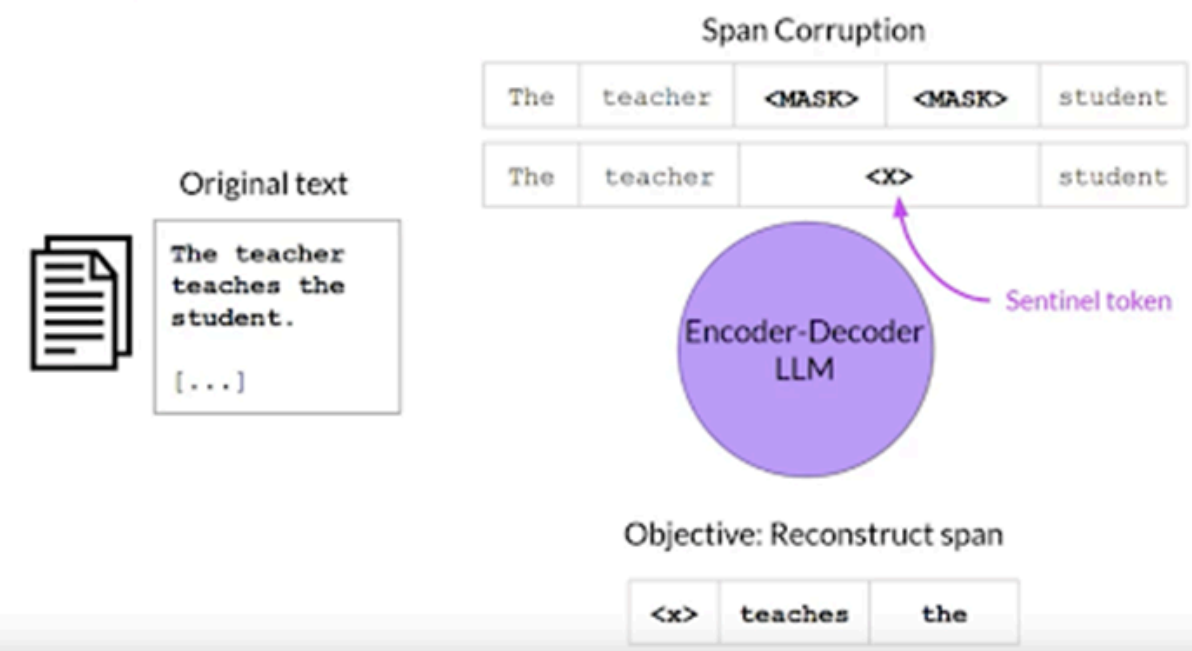
Training objective is to predict the next token

Pre-training of Encoder-only LLM



Training objective is to reconstruct the masked text

Pre-training of Encoder-Decoder LLM



Training objective is to reconstruct the span of tokens called the sentinel token

PRE TRAINING OF LLM'S

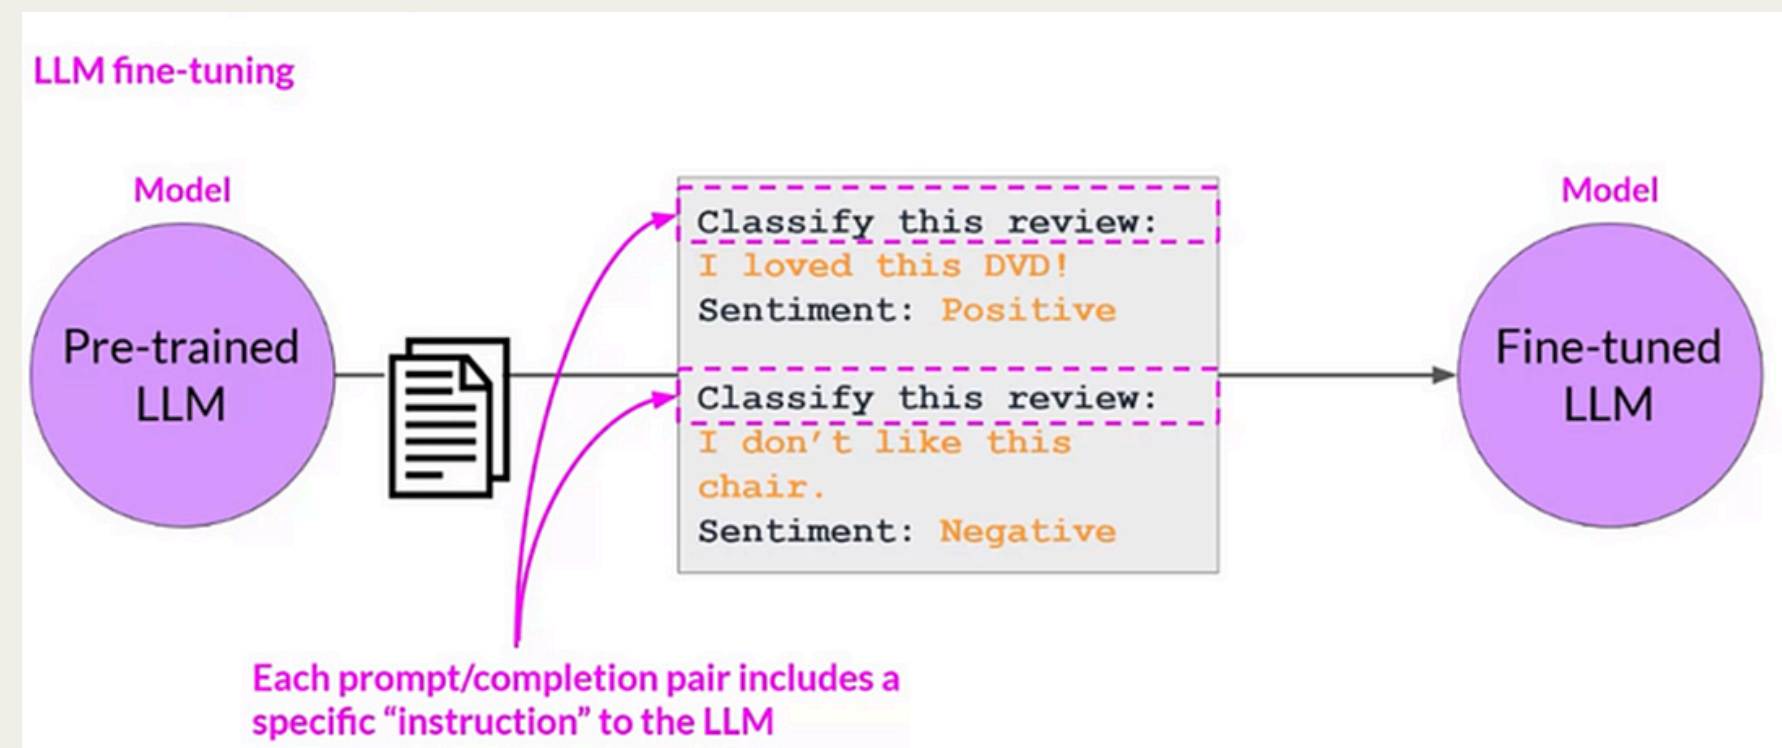
- Steps of Pre-training:
 - Training Process
 - Model starts with random weights and learns patterns from the data
 - Processes text in chunks/windows (typically 512–2048 tokens)
 - Updates its parameters using backpropagation and optimization algorithms
 - Training continues until performance plateaus or computational budget is reached
 - Outcomes
 - Model develops broad understanding of language patterns
 - Learns grammar, facts, reasoning capabilities
 - Acquires general knowledge across many domains
 - Can perform basic tasks without additional training

FINE TUNING OF LLM'S

- In context learning may not work for smaller models since examples take up space in the context window. Instead try fine tuning the model
- Fine tuning is taken up supervised learning with labeled examples.
- Aim of fine tuning: Updating the weights of LLM to gain from space and increase accuracy of the model

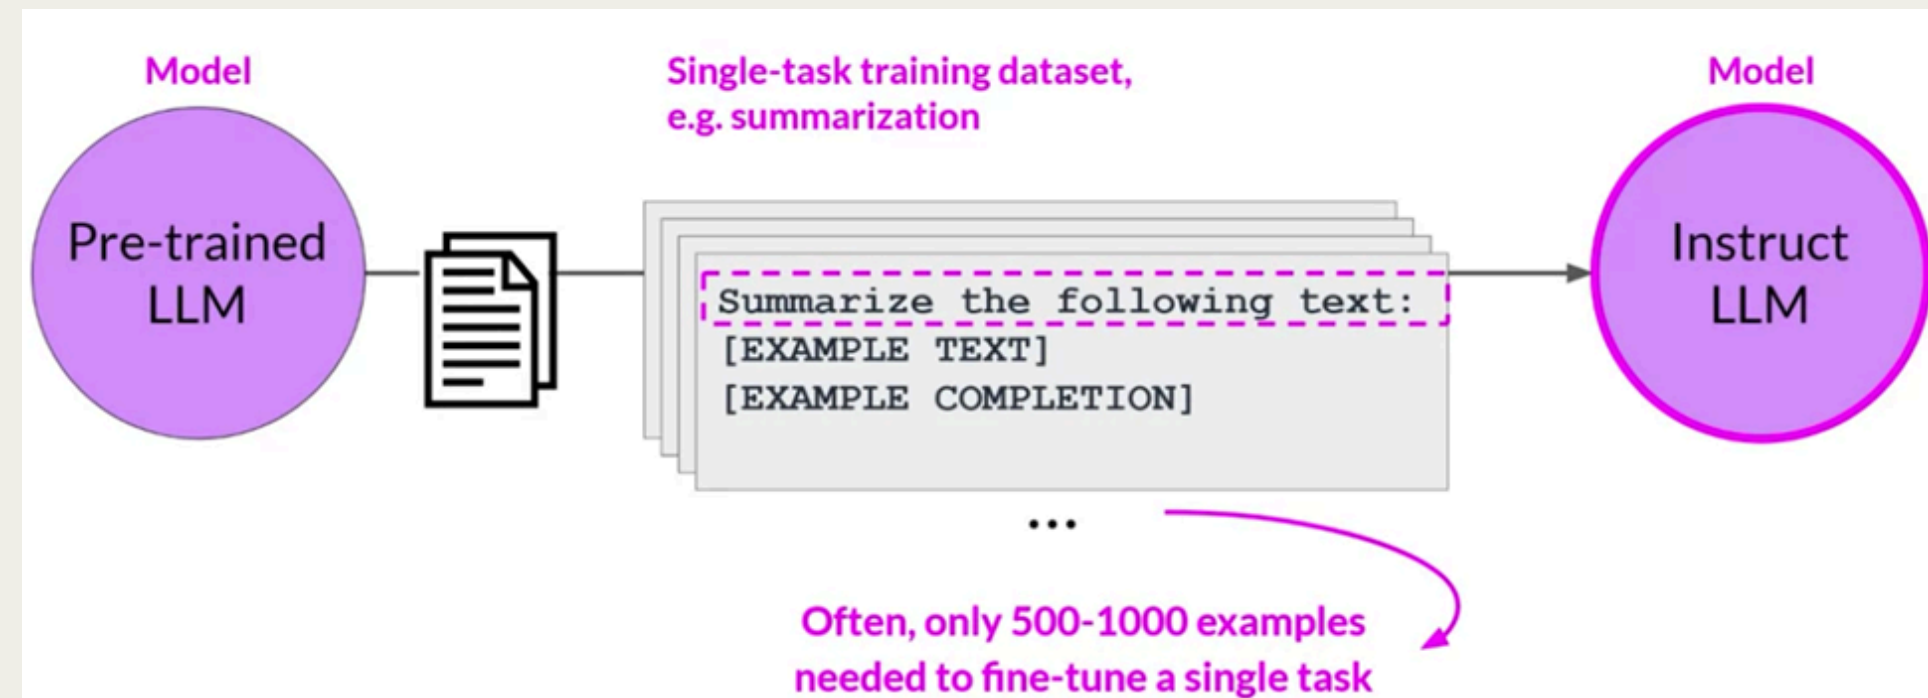
Instruction Fine Tuning:

- Good for improving a model's performance on a variety of tasks,
- IFT trains the model using examples that demonstrate how it should respond to a specific instruction



FINE TUNING OF LLM'S

- Fine tune a pretrained model to improve performance on only one task
- 500 – 1000 examples is enough



- Potential downside:
 - Catastrophic forgetting:
 - Models tend to forget what they were originally trained on when you do a fine tuning. If you do fine tuning too much, you may end up with catastrophic forgetting

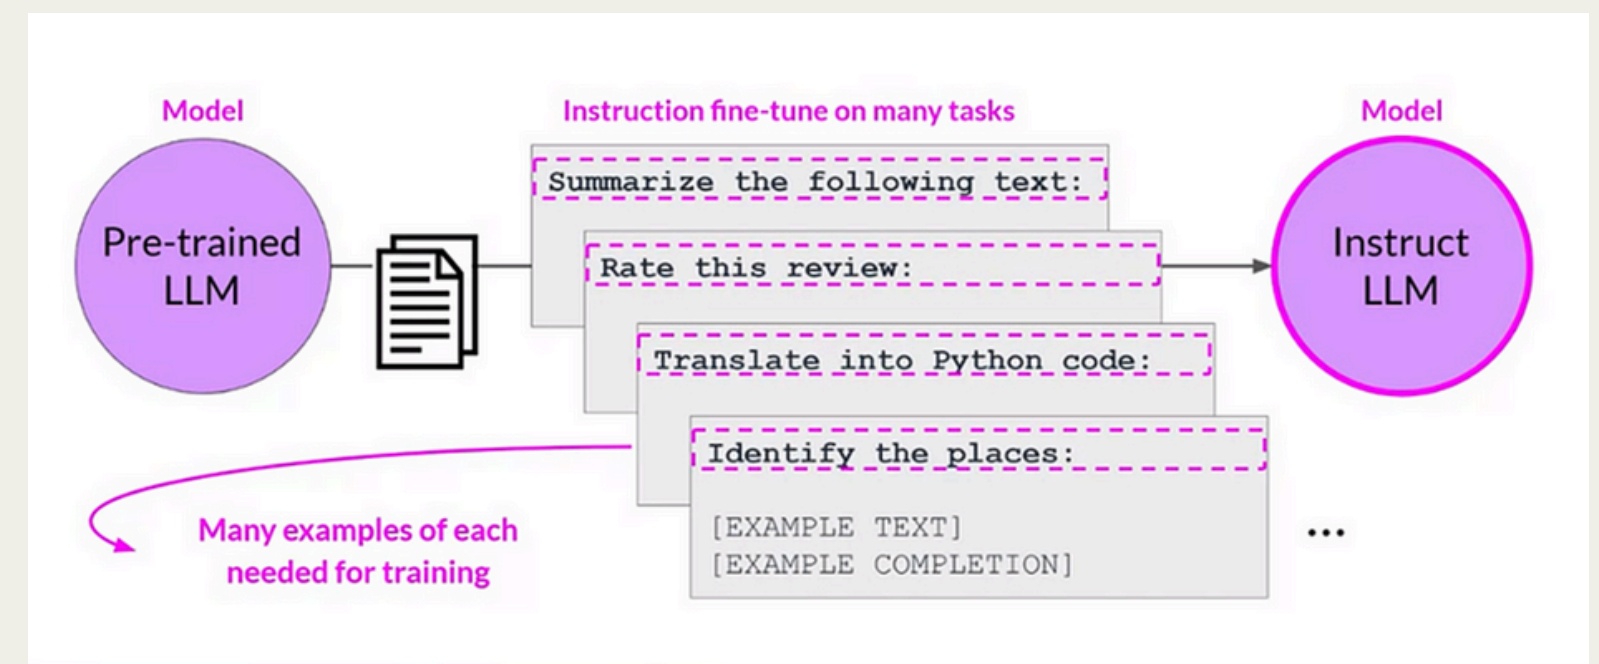
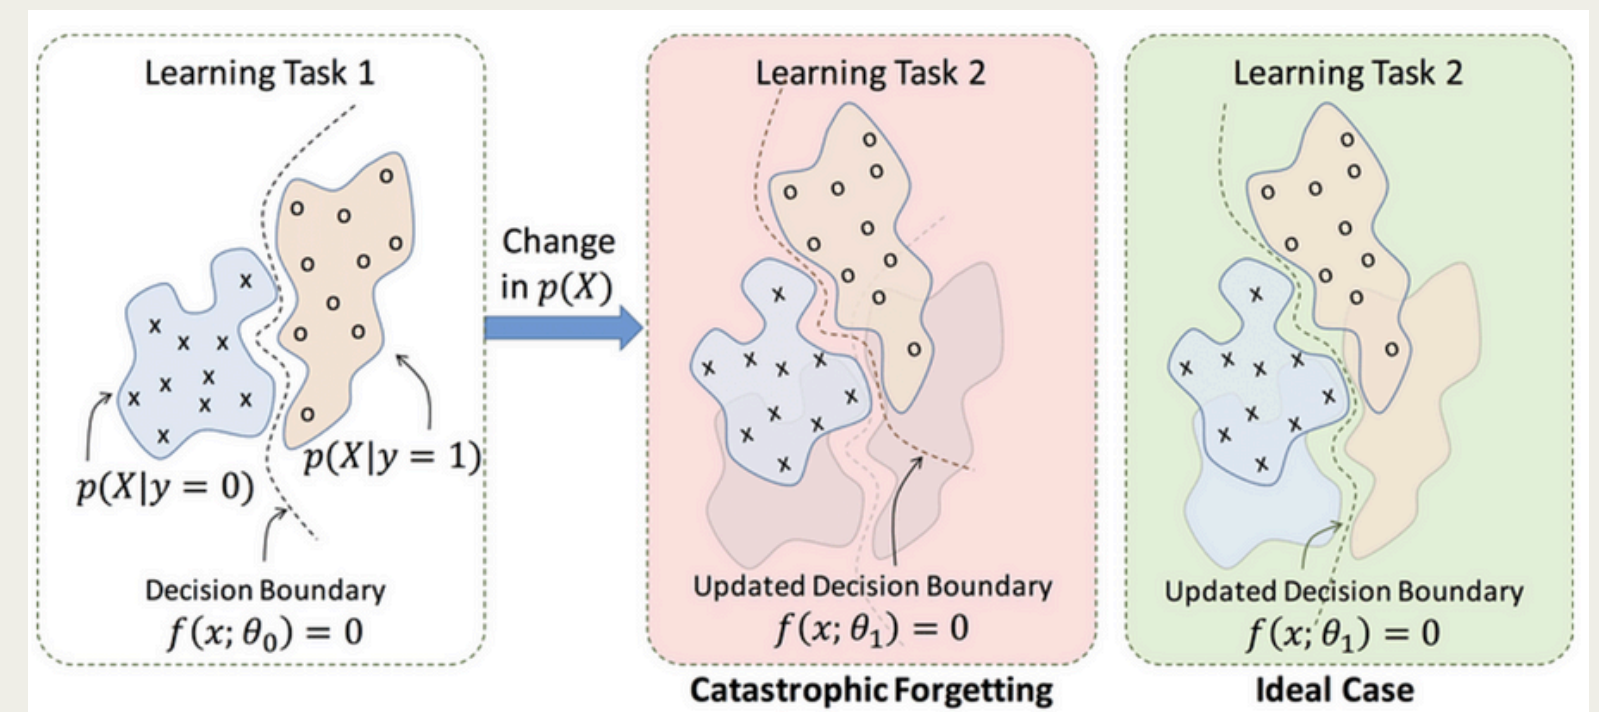
FINE TUNING OF LLM'S

- How to avoid catastrophic forgetting?
 - 1) Do not have to
 - 2) Fine tune on multiple tasks at the same time
 - 3) Consider parameter efficient fine tuning (PEFT)

- Multi-task instruction fine tuning:
Training dataset is comprised of example inputs and outputs for multiple tasks

Drawback:

Requires a lot of data 50 – 100000 examples in training data set




FINE TUNING OF LLM'S

- Steps of Fine Tuning:
 - Loading and preparing the dataset
 - Preprocessing the data
 - Model preparation
 - Initial Predictions
 - Training the model
 - Evaluating the model
 - Making predictions with the trained model

FINE TUNING OF LLM'S


- PEFT
- Freeze the original model weights.
Add a small number of new Parameters or layers and fine tune only new components



- PEFT can be performed on a single GPU
- PEFT is less prone to catastrophic forgetting problems of fine tuning

FINE TUNING OF LLM'S

- LoRA
- Freeze the original model's weights
- Inject 2 rank decomposition matrices (adapters)
- Train the weights of the smaller matrices



THE DIFFERENCE BETWEEN PRE-TRAINING AND FINE-TUNING

- **Pre-training:**

- Uses vast amounts of general text data
- Higher learning rates
- More computational resources
- Builds foundational capabilities
- More expensive and time-consuming

- **Fine-tuning:**

- Uses targeted, task-specific data
- Lower learning rates
- Less computational intensity
- Preserves general knowledge while specializing
- More accessible to organizations with limited resources

TOKENIZATION ON LLM'S

- Tokenization is the process of breaking text into smaller units called tokens, which can be words, subwords, or characters. LLMs process text as tokens rather than raw words.
- How Tokenization works:
 - Word-based: Splits text into words (e.g., "Hello world!" ➞ ["Hello", "world!"]).
 - Subword-based (Byte-Pair Encoding, WordPiece): Splits words into meaningful parts (e.g., "unhappiness" ➞ ["un", "happiness"]).
 - Character-based: Each character is a token, useful for handling unknown words.
 - Byte-level (GPT-style): Works at the byte level, handling all languages and special characters efficiently.
- Good tokenization improves model comprehension and generation quality.
- Proper tokenization directly impacts model accuracy, speed, and cost in LLMs.

EMBEDDINGS ON LLM'S

- Embeddings are numerical vector representations of words, phrases, or sentences that capture their meanings in a multi-dimensional space. LLMs use embeddings to understand and process language efficiently.
- How embeddings work:
 - Conversion: Words/sentences are mapped to dense vectors (arrays of numbers).
 - Semantic Relationships: Similar words have closer vectors (e.g., "king" and "queen" are near each other).
 - Context Awareness: Modern embeddings (e.g., from transformers like BERT) capture meaning based on context. Same word can have different vectors based on usage
- Types of Embedding:
 - Word2Vec, GloVe: Static word embeddings (fixed meaning per word).
 - Transformer-Based (BERT, GPT): Contextual embeddings that change based on sentence structure

EMBEDDINGS ON LLM'S

- Modern embeddings typically use 256–1024 dimensions
- Each dimension may capture different aspects of meaning
- Higher dimensions allow for richer representations
- Trade-off between expressiveness and computational cost
- Words with similar meanings cluster together in vector space
- Captures multiple types of relationships:
 - Synonyms (happy/joyful)
 - Antonyms (hot/cold)
 - Analogies (Paris:France :: Berlin:Germany)
 - Categories (apple/orange/banana)

EVALUATION OF LLM'S

- **Perplexity (PPL):** Measures how well a model predicts the next word in a sequence; lower values indicate better performance. Calculated as the exponential of the average negative log-likelihood. Useful for comparing models during training
 - Limitation: Doesn't directly measure task performance or output quality
- **Accuracy:** Percentage of correct predictions
- **F1 Score:** Balances precision and recall
- **ROC-AUC:** Measures discrimination ability
- **BLEU:** Compares generated text with reference translations
- **ROUGE:** Measures overlap with reference summaries
- **METEOR:** Considers synonyms and stems in evaluation
- **GLUE/SuperGLUE:** Natural language understanding tasks
- **Humans rate** outputs for: Quality, Coherence, Factual accuracy, Helpfulness
- **A/B Testing**
 - Compare outputs from different models
 - Humans choose preferred response

REAL WORLD APPLICATIONS OF LLM'S

- Chatbots & Virtual Assistants – Customer support, healthcare Q&A, AI tutors (e.g., ChatGPT, Alexa).
- Content Generation – Marketing copy, blog writing, code assistance (e.g., GitHub Copilot).
- Translation & Language Processing – Real-time translation, transcription, localization.
- Education – AI tutors, automated grading, personalized learning.
- Healthcare – Medical documentation, drug discovery, diagnosis assistance.
- Finance & Business – Fraud detection, market analysis, legal document review.
- Personalization – Product recommendations, content curation, HR screening.
- Scientific Research – Paper analysis, AI-driven simulations, debugging assistance.

Thank you!

Gülşen Sabak

Bogazici University Computer Engineering